Enterprise's Last Exam (ELE)

Model Comparison Report

Benchmark of organizational-reasoning ability across 161 enterprise scenarios (156 community submissions + 5 reference scenarios). Each scenario tests whether a model can reason over fragmented, contradictory enterprise context (CRM, Slack, email, policy documents, billing systems) and reach the decision an experienced operator would make.

Judge: GPT-4o-mini (reference-based, fixed across all models). Date: July 2026.

Overall Results

Rank	Model	Access	Accuracy	Correct	Avg Latency
1	Kimi K2.5	Bedrock	87.0%	140/161	1443 ms
2	Claude Opus 5	Bedrock	85.7%	138/161	3853 ms
3	Claude Sonnet 5	Bedrock	84.5%	136/161	4488 ms
4	GPT-4o-mini	OpenAI API	77.0%	124/161	847 ms

Accuracy by Difficulty

Difficulty	Kimi K2.5	Claude Opus 5	Claude Sonnet 5	GPT-4o-mini
Hard	80%	89%	84%	80%
Expert	93%	82%	85%	74%

Accuracy by Reasoning Category

Category	Kimi K2.5	Claude Opus 5	Claude Sonnet 5	GPT-4o-mini
Approval Chain	93%	96%	96%	86%
Cross System Synthesis	90%	83%	80%	70%
Entity Resolution	90%	93%	90%	90%
Policy Version	88%	81%	84%	75%
Precedent Exception	76%	76%	76%	65%
Temporal Consistency	79%	79%	75%	71%

Key Findings

• **Kimi K2.5 leads overall (87.0%)**, narrowly ahead of Claude Opus 5 (85.7%) and Claude Sonnet 5 (84.5%). GPT-4o-mini trails at 77.0%.

• **Kimi is strongest on the hardest tier** — 93% on expert scenarios, well above the Claude models (82–85%). It also leads cross-system synthesis (90%).

• **Claude Opus 5 leads on approval-chain and entity-resolution** reasoning (96% / 93%), reflecting strong structured-authorization logic.

• **Precedent-based exception handling is the universal weak spot** (76% for the top three, 65% for GPT-4o-mini) — the category requiring judgment about when to bend a rule.

• **Latency vs accuracy:** Kimi delivered top accuracy at ~1.4s/scenario, while the Claude models were 3–4x slower. GPT-4o-mini was fastest but least accurate.

Methodology

Scoring pipeline: exact match → LLM-as-a-judge (for non-exact answers). Multiple-choice answers are scored by exact letter match (markdown formatting such as **A** is normalized before extraction). The judge (GPT-4o-mini) grades free-text answers against a held-out correct answer and rationale that are never shown to the model under test. All models received identical prompts with no hints. A scenario counts as correct when its final score is ≥ 0.5.

Note: an earlier scoring pass under-counted Claude Opus 5 because it wrapped multiple-choice answers in markdown bold; the extractor has been corrected and all results re-scored uniformly.